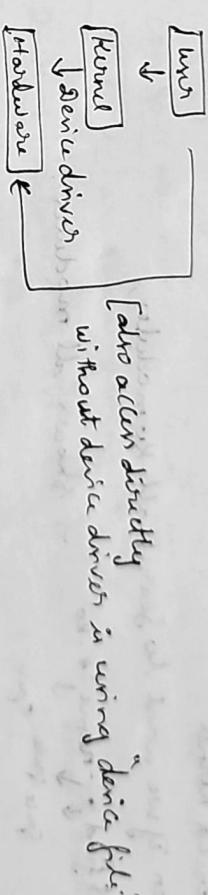


① Device driver:

→ It is a piece of software that controls a particular type of device which is connected to the computer system.

3 sides: → ① one side talks to rest of kernel.

- ② one talks to the hardware
- ③ one talks to user.



② Kernel module:

Traditional way of adding code to kernel was to recompile the kernel and reboot the system.

→ Kernel modules are piece of code that can be loaded/injected and unloaded/removed from the kernel as per demand/need.

other names of kernel modules:

1. Loadable kernel module (LKM) 2. Modules

Extension: .ko (kernel object)

Standard location:

modules are installed in /lib/modules / kernel version > directory of soft by default.

Kernel

cd /lib/modules

5.1 generic

To check the current kernel

\$ uname -r

5.4.0-1-generic

then cd kernel

then if we want to search all kernel modules,

\$ find . -name '*.ko' → shows all modules

space space space

Count:

9 find . -name '*.ko' | wc -l

12 5505

(3)

Kernel module vs Device driver

→ A kernel module may not be a device driver at all.

→ A driver is like a sub-class of module. Device drivers

→ A driver is always a K.M.

modules are used for below:-

1. D.D 1. Network drivers: Drivers implementing a network stack

2. file system 2. Network protocols (TCP/IP)

3. system calls 3. Terminal devices

System flavor → always in ram.

First (Bare Kernel) is loaded (booted), after it is booted Kernels are loaded. K.M may be in separate memory.

Advantages of K.M:

1. All parts of Bare Kernel stay loaded all the time. Modules can save you memory, because you have to have them only when you are actually using them.
2. Users would need to rebuild and reboot the kernel every time they would require a new functionality.

3. A bug in driver which is compiled as part of kernel will stop system from loading, whereas module allows systems to load.

4. Easier to maintain & debug multiple systems.
5. Makes it easier to maintain machine on single kernel box.

Disadvantages:

1. Size: Module management consumes unpageable kernel memory than an equivalent kernel with drivers compiled into kernel image itself.
- A basic kernel with no of modules loaded will consume more memory than an equivalent kernel with drivers compiled into kernel image itself.

This can be very significant issues on machines with limited physical memory.

2. As the kernel modules are loaded very late in the boot process, hence core functionality has to go in bare kernel (e.g. Memory Management).

3. Security: If you build your kernel statically and disable Linux's dynamic module loading feature, you prevent run-time modification of kernel code.

Configuration:
In order to support modules, kernel must have been built with following options enabled:

CONFIG_MODULES=y
CONFIG_MODULE_UNLOAD=y
CONFIG_MODULE_FORCE_UNLOAD=y
CONFIG_MODULE_UNLOAD=y
CONFIG_MODULE_FORCE_UNLOAD=y

To check that:

\$ cat /boot/config-`uname -r` | grep CONFIG_MODULES

\$ cat config-`uname -r` | grep CONFIG_MODULE_UNLOAD=y

Kernel modules are loaded via /etc/modules or /etc/modprobe.d/*.conf

or /etc/init.d/*

\$ cat /etc/modules

[Output]

CONFIG_MODULES=1	CONFIG_MODULE_UNLOAD=y	CONFIG_MODULE_FORCE_UNLOAD=y
------------------	------------------------	------------------------------

④ Types of Modules:

1. In-source tree: Modules present in the Linux kernel source code. [Linux source code → kernel.org]
2. Out-of-tree: Modules not present in Linux kernel source code.

All modules start out as "out-of-tree" development, that can be compiled using context of source-tree. Then went to review. Once a module gets accepted to be included, it becomes an in-tree module.

⑤ Basic Commands:

1. List modules:

(lsmod) lsmod gets its information by reading the file /sys/module. Currently the module (proc/module) previously registered.

lsmod

cd /boot

\$ lsmod

It shows all modules

lsmod | less \Rightarrow Shows latest installed module

Module	Size	Used by
ccm	20482	(in) dependency
intel_rapids	24546	intel_rapids
intel_rapl_common	0	

2. Module Information:

(modinfo): prints the information of the module

modinfo \rightarrow Module name is

intel_rapl

It shows info.

Chosen from lsmod (Previous Command)

eg:

modinfo intel_rapl | less, prints at the bottom (bottom)

It shows info. [Please check it in terminal]

One interesting, we can find whether

it is in-tree or out-of-tree.

"intree: Y" \rightarrow will appear there.]

⑥ Hello world kernel module-

In C/C++ programming, we have main() as entry point & exit point.

Kernel modules must have atleast two functions:

\rightarrow ① a "start" (initialization) function: which is called when the module is loaded into kernel.

\rightarrow ② an "end" (cleanup) function called: which is called just before it is removed.

This is done with module_init() & module_exit() macros.

Linux: Module should specify which license you are

using MODULE_LICENSE macro

④ GPL [GNU public license v2 or later]

Header files:

Every kernel module needs to include linux/module.h for Macro expansion of module_init & module_exit linux/kernel.h only for macro expansion for the printf() log level.

steps in code: how to make this module
① include header file

② ~~include license~~ only when we want to use module

③ include initialization & exit fun-

vi world.c

#include <linux/module.h>

In order to build kernel module, we need to use kernel make file which is in path /lib/modules/`uname -r`/build/Makefile

to create kernel module which we are going to build

then we can't directly use kernel makefile, because we have to include our object in that file (i.e) world.o

so, we are creating a Makefile in local path & use kernel to check this created makefile.

vi Makefile → caps.H

obj-m := world.o

Two Options:

① -m => It tells where building modules

② -y => It builds kernel module code

③ To Build Module:

make -C /lib/modules/`uname -r`/build M=world.o

check in local folder => It creates .ko file & others (i.e world.ko)

To check whether module is inserted or not,

\$ lsmod | less

↳ Shows the module list where we can see world.ko

But it will not be available - because we are not inserted into kernel. Let's do

① Invoking KLD:

\$ sudo insmod ./helloworld.ko
Now check lsmod | less, we can see that.

② Now check lsmod | less, we can see that.

③ And also in /sys/module/ folder is also created.

④ \$ dmesg (or \$ sudo dmesg) \Rightarrow log file is also created where we can see print statements module entry will be there.

⑤ Remove the module.

\$ sudo rmmod ./helloworld.ko

⑥ Clean the module [Clean the files, folder related to KLD]

\$ make -C ./helloworld module/uname -r /build/mkdir} clean.

⑦ printf vs printk:

printf() is a function in standard library.
printk() is a kernel level function. (task_struct->printk)
printk() is called with one more argument than printf().
Like below

```
printf("HelloWorld\n");
```

and

```
printk(KERN_LOG_PRIORITY "HelloWorld\n");
```

Here LOG_PRIORITY is one of eight values (prefixed in kernel/include similar to /usr/include/sys/syslog.h)

EMERG, ALERT, CRIT, ERR, WARNING, NOTICE, INFO, DEBUG

Priority decreasing in this order.

DEBON

⑧ printk() writes to the kernel buffer, whereas printf() writes on the standard output.

⑨ Simplified Makefile:

\$ man make

make [options] [target]

(basically) takes care of all tasks which are in makefile
-C

Instead of giving whole command make -C /helloworld we can simply give

make to kernel it takes care of kernel file auto linking and building of modules if needed.

① \$ make

② \$ make clean

Let's do

\$ vi Makefile

Obj-m := World.o

all: > target

make -C /lib/modules/`uname -r`/build M=\$PWD

clean: target

make -C /lib/modules/`uname -r`/build M=\$PWD clean

modules

as it calls init-module() to intimate the kernel that a module is attempted to be loaded & transfer the control to kernel.

Kernel is running in super privileged user, the only way to allow the kernel by unmapce program in system call.

b) In kernel, op_init-module() is run. It does a sequence of operations as follows:

→ verifies if the user who attempts to load the module has the permission to do so (or) not

→ After the verification, load module fn is called.

→ The load module function assigns temporary memory

and copies the elf module (our module format) from user space to kernel memory using copy-from-user-fn.

→ It then checks sanity of ELF file (verification if it is a proper ELF file).

→ Then based on ELF file interpretation, it generates offset in temporary memory space allocated. This is called Convenience Variables.

→ user arguments to the module are also copied

as it calls init-module() to intimate the kernel that a module is attempted to be loaded & transfer the control to kernel.

Kernel is running in super privileged user, the only way to allow the kernel by unmapce program in system call.

b) In kernel, op_init-module() is run. It does a sequence of operations as follows:

→ verifies if the user who attempts to load the module has the permission to do so (or) not

→ After the verification, load module fn is called.

→ The load module function assigns temporary memory

and copies the elf module (our module format) from user space to kernel memory using copy-from-user-fn.

→ It then checks sanity of ELF file (verification if it is a proper ELF file).

→ Then based on ELF file interpretation, it generates offset in temporary memory space allocated. This is called Convenience Variables.

→ user arguments to the module are also copied

But check kernel buffer by "dump", there are many statements
"Trunk allocation". But after it errors.

to kind memory
is done.

\rightarrow system resolution - it can return a reference.

→ load - module ,
→ signed by load - module

Know more about behavior. Continue to moderate. Mind limit that has a last.

β is added to doubly linked lists in the system keep to groups.

of all modules loaded from the kernel module

⑩ what happens if we run too many observations with one

in word spaces:
in a delora book, nothing is ever written
in more than one space.

\rightarrow This means, multiple independent failures in parallel systems

return of returning to a place where

But in bound space? precise gravitation
is not possible?

static int testInit(void)
{
 /* ... */
 Info("S: Initialization in -function");
 /* ... */
}

printek (REK-...)

3
and all the rest of us were present in the hall.

trying to build it wrong i.e. trying to invert module \rightarrow error "permissions" issued.

(giving another name to module)

g. I have hello.c that contain all of my local code.

+ But I want models to be linear. No.

↳ vi Makfile
=> modified
↳ vi Makfile

obj - m := hello.0

↳ [func - obj] := hello-o

St. L. Y. 1000.0

For Xmas, I have a
first it generates hello.o → then linus.o → linus.hello

⑫ Kernel Module kann aus mehreren Dateien bestehen

W.R.C - Job function

```
#include <linux/kernel.h>
```

```
void func(void);
```

```
print("Hello World")
```

111

modification in makefile: (file wx.c)

in
previo

We will lose the boot up logs if we don't store them in the buffer.

~~obj - mi :=~~ ~~word~~
world - objs := word func.0

Important note from Handel's office.

Important note: Before cleaning the models, must remove model

before naming.

modification in Makefile

obj-m:=hello.o

Obj -mf = hello2.o

all.

clean.

નીચેની લાંબી વિશ્વાસી કાર્યક્રમની પ્રાણી અધ્યાત્મિક વિદ્યા

14 What does a

skewer keeps an

Systog daemon si

in / van / kog / drenq.

dmrg Command is used to control (or) print kernel ring buffer. Dfault is to prints messages from the kernel ring buffer or to console.

460 / van Hoorn

gmag (or) f Sudo dmag

Important dmcg Commands

1. Clean Ring buffer:

→ will clear the ring buffer after printing

String - C will clear the ring buffer but does not points

on the course. The second day was spent in the same way.

2. Don't point Timestamps! Just as I said, get rid of timestamps.

\$down-t → will not print timestamps

and the first time I saw him he was wearing a
blue shirt.

Previously I have written

01/14/32 781953] Ent - Hello 2 - east - In east

after drying - t

How about new XT-Cardboard, and so on
for paper? - This will be
done.

For - women - love : in life go with whatever comes

③ Put all drug commands to list of hooks :-

Ex: drug -l err
w [2.93894] qmpolicy: failed to initialize policy for policy
with error

Ex: drug -l err, warn

op: +

[5.47667] system-journald[468]: File /var/log/journal/

System (or) unclearly shut down.

Print human readable timestamp

\$ drug -t will print timestamp in readable format. Drugs b

driving -t will print timestamp in readable format. Drugs b

Note: timestamp could be inaccurate. Drugs act as

⑤ Display the log level in the output instead of just good

driving -x will add loglevel to the options < t -g -w -f

log level

⑥ we can combine options, no driving -tx will print both

human readable time & log level. It does. Consider that

④ Drug follow option:

Previously what are we doing? Load the module → print the kernel msg (using drug) & unload the module → print the kernel msg.

Now, no need to give drug after ins & mod. Now, we are going to run drug in background.

\$ drug -w &

op (log comes)

\$ jobs

[1]+ 0+ running

driving -w

\$ sudo insmod ./hello.ko

[2]+ 67+ running

tut(hello).drivenit → immediately

printing without exit function → giving drug

⑦ without exit function?

Same code without exit fn.

→ tried, code is compiled into hex output with a character like
→ normal → works → tut - init → printed

→ command work?

Sudo wood hollow

remove -

8 cd lined - 5.2.81

✓ cd lined - 5.2.81

Why do we get - EBV when we remove the module?

24

Definition: Kernel/module - c
→ there is a check in implementation to strictly have
the exit function if there is a init function. If it failed
to find exit function it throws - EBUSY

* * * * *

if (mod < init || mod > max), and also if mod is odd, because then it is not possible to split mod into two equal parts.

Δ fixed = $\Delta_{\text{true}} - \Delta_{\text{false}} = \text{unloaded} (\text{flexure})$ [51 pp. (c) 255]

if (locked) {

忙 = Busy

3 got out;

3

dele- module in the system call which implements the `lprm` -

removing/unloading the module.

1. Writing of Compiling Kernel modules

③ understanding Modules, symlinks and modules-order
modules-order: In case you are Compiling multiple modules together, it will Lost out the order in which Compilation and creation of .ko takes.

`lsmod` → List modules that loaded already
`insmod` → Insert Module into kernel

二十一

remove → Remove module from kernel
modprobe → Add (or) Remove module from kernel

~~differentiate L/H in normal vs hypothyroid~~

modprobe: Load the module given in `modpath/to/module.so`

nedi medprobe/kompl/tut/hello ko will not work.

invmod: Dependencies if present are not loaded.
modprobe: modprobe calculates dependencies, loads the module.

dependencies and then the main model

that is defined in your module and will not open in the modules you see of kernel.

(2) How would be Callisto's dependence
on Earth tool to Gal

Madprob depends on degraded knowl.

Dependents calculate - listing & replace the dependency

1.1b) module/(unname -s) file , 1.1b) (unname -s)/module .dip file.

information in //lib/procmail/

we have now begun to work.

...Kernell - reprezentácia
...výberu výrobkov

1000 Washington Avenue, C. P. O. Box 2111, Toledo, Ohio.

Stay when you say good...
I'll stand then adm & 41-ko.

loaded first and then
lifted right to left & removed left to right.

Modules are loaded in the order of module = exit function

(2) understanding module init & module = ...

We know or inferred, the function paroxysm in *measles* does not occur in all cases.

macro is called, and on demand, generates function definitions for its arguments.

module unit called int which contains method getoption

its see definition of this macro, it is present in `Linux/modulen.h`.

that's why we pass - history - enough to make

If you declare module_init function which returns void instead of int, the compiler will throw warning.

initial_t is defined in Linux/init.h
typedef int (*initial_t)(void);

during compile time, the function passed to macro is compatible with initial -t type-

void cleanupModule(void) — automatically calls (#define)

```
Static inline initcall_t - outlist(void)  
{  
    return exitfn; }  
  
#include <linux/init.h>
```

```
#define module_exit(exitfn) \
```

```
{ return initfn; }
```

```
#define module_init(initfn) \
```

Twenty four
inches

Static inittry fun(void) {
 // module - static (inittry)

Kernel programming also allows us to pass command line parameters.

Command line parameters provide a single Linux driver to do multiple things, for example

- Instead of fixing to single I/O address for read/write, it can provide that as command line argument and allow user to read/write any address.
- Enable/disable debug log/printk
- Allow user to set the mode if driver supports multiple modes.

How to pass parameters to module?

We can add parameters using module-param macro, declared in module-param.h file.

#define module_param(name, type, perm)
module_param(name, type, perm)

name: name of variable which will be returned
type: type of variable supported types are char, p, bool, int, long, short, ushort, ulong, ushort

perm permission for the sysfs entry.

Ex: S-I-RWNO: Only read by all users

O: No sysfs entry.

You can also use numeric values like 0644 for permission entry.

Example: watch video.

Module param Example:

```
#include <linux/module.h>
#include <linux/moduleparam.h>
```

```
int basic=0;
```

```
int hex=0;
```

```
int allow=0;
```

```
char *int Salary_init(void)
```

```
{
```

```
    printk(KERN_INFO, "I'd", basic+hex+allow);
```

```
    return 0;
```

```
static void Salary_exit(void)
```

```
{
```

```
    printk(KERN_INFO, "exit");
```

3. When we run the module, we get the output as

```
module_init(Salary_init);
```

```
module_exit(Salary_exit);
```

```
-> exit
```

Parsing Parameters

\$ sudo insmod ./arg.o banic=18000 loop=3000 allow=1000

if Take home = 20000 [Note previously we have to do \$ sudo insmod -v arg.o to see this as soon as]

29. What happens if we pass incorrect value to module parameter?

\$ sudo insmod ./arg.o banic=0.005

\$ sudo insmod ./arg.o banic=0.005
Error: Could not insert module ./arg.ko
ignoring: Error: Could not insert module ./arg.ko
Module ./arg.ko has invalid parameters

(bio) (info) (in note) (info)

we can also see error in dumping module

30. How to pass parameters to built-in modules?

How to find out values of already loaded modules

\$ cat /sys/module/module-name/parameters/parameter

This is only possible, if the permission field in the module parameter file is not zero.

(dmesg) ftrace -d 1000

How can we pass arguments to modules which are built-in

Module parameters for built-in modules are passed through kernel command line.

Syntax: <module-name>.<parameter-name>=value

So if the argument module was built-in, we can append

to the kernel command line for passing 18000 as banic

How can we pass arguments which are called by modprobe?

modprobe reads /etc/modprobe.conf file for parameters

31. How to pass string with multiple word as parameter?

Now we have character variable (name) in our module.

\$ sudo insmod ./arg.o name="Hello world"

dmesg -r

World is ignored (because) when (and) the

key of sed " was too long = HelloWorld" was used

\$ dmesg -r
Hello world was used to copy from user, but

'world' is ignored

try:
 name = "Hello world"
 \$ sudo ./name

gives
Hello world.

giving → ~~off~~: Hello world.

Reason:

This happens because shell removes double quotes and
pars it to instead, to avoid this add a single quotes over the

string.

Run the following command: "inverted argument". It does not

name = "Linux world" to pass the whole string as one

3. Parsing zero to permission argument of module

- param macro

character at will.

Code changes in exec.c: -(return code), pointer

module-param(name, charp, 0); parsing zero

module-param(name, charp, 0); boom! crash

When we go & check,

#ls /sys/module/module/parameters

shows only one parameter: (0) loop count

But "name" is not present because we pass zero

to permission argument

But when we check module info, [because it can access
the module's arguments] it shows here but not there

param: (name:char)
param: (loop-count:int)

(3) Parsing array as module parameter:

To parse multiple parameters, we need to pass parameter
array.

To parse array we need to use module-param-array()

function instead of module-param() function

#define module_param_array(name, type, num, perm)

Here num: optional pointer filled in with number

written.

\$ vi array.c

#include "moduleparam.h"

#include <linux/moduleparam.h>

int array[4];

Static int arg_count=0;

" " kind, it will be present

when you run /boot/system.map -l (linux -junk) in

I try to eat, however.

+ Et will lust hot. "Kobjins-ops-
f so, g yado en -

③ Exporting symbols: Symbols?

How to export your J

3 when you define a new function ...
when you define a new function in local, only the
default behavior of this function will change.
the place in which the function is defined can call it, cannot be

accessed by other modules.

To export this module, we need to write a file named `symbolopt.relops` in `libdwarf` directory. It contains the following code:

```
(or) EXPORT_SYMBOL(opt_relops);
```

Once you export them, they will be available to other modules. You can't use them directly without opening kernel header files.

Elegans Lepidoptera

Export symbol: The exported symbol can be used by any kernel module.

EXPORT-SYMBOL-GRPL: The exported symbol can be used by only local libraries.

→ Jupyter notebook code.

/proc/kallsyms: Contains symbols of dynamically loaded

modules as well as built-in modules.

... contains symbols of every built-in module. Let's say we write a module & export its symbol. It will

be available in /proc/kallsyms only. Because what we

38. Linux Kernel module example of exporting function

```
#include <linux/list.h>
```

MOBILE-LICENSE ("Copy") wird als private Lizenz für den Verwendungsbereich des Nutzers bestimmt.

1900-1901 (Vol. 1) reading us some interesting books.

printk(KERN_INFO): g: jffs2:initial, -func, if
-- (bootlog) evntk:0x0000000000000000

make int -t name so loader knows what to do

static void

EXPORT_SYMBOL_C(fprint_jiffies)

do_brenda(jiffies)

module_init()

module_exit()

own linker of example:

4. write module which exports a function myadd performing addition of two numbers:

mod1.c

```
#include <linux/module.h>
int myadd(int a, int b)
{
    return a+b;
}
```

prefix ("mod1: Adding %d with %d result :%d\n", func, a, b, add);

3. write module which has symbol to add to loader binary

EXPORT_SYMBOL(myadd);

static int module_init(void)

module_init(myadd);

Example of module stacking in Linux kernel

① Msdos filesystem relies on symbol exported by jffs module

② Parallel port printer driver (lp) relies on symbols exported by generic parallel port driver (parport)

3. static void module_init(void)

3. static void module_exit(void)

module_init()

(mod1) depends on mod2

mod2.c

```
#include <stdio.h>
int myadd (int a, int b);
static int myadd (int a, int b)
{
    return a+b;
}
```

```
printf("1: add: %d\n", -fun c, myadd (3,5));
return 0;
```

```
(char *a, char *b) {
    int i;
    for (i=0; a[i] != '\0' && b[i] != '\0'; i++)
        a[i] = a[i] + b[i];
    a[i] = '\0';
}
```

```
int myadd (int a, int b)
{
    return a+b;
}
```

① Create makefile which builds the both mod1 & mod2.c files

② Create mod1

③ Sudo cat /proc/kallsyms | grep myadd

④ Invert mod2

⑤ dmesg → result will appear

40 · Linux kernel module sample of exposing variable

⑥ symbol-expert.c:

⑦ #include <...>

#include "symbol-test.h"

Module-Name ...

Struct test foo;

EXPORT_SYMBOL(foo);

symbol-test.c

```
static int test_export_init(void)
{
    foo.a = 5; foo.b = 6; foo.c = 4. ...
    return 0;
}
```

```
static void test_export_exit()
{
    /* do cleanup */
}
```

```
symbol-test.h
```

```
#ifndef _SYMBOL_H
#define _SYMBOL_H
```

```
struct test
{
    int a; int b; int c;
};
```

```
extern struct test symbol_test; // no file symbol-test.h
```

① Invert symbol-expert.c → kernel module registered at my kallsyms

② Sudo cat /proc/kallsyms | grep foo

③ vi & symbol-test.h → kernel module found in /sys/module/symbol-test/

symbol-test.c

#include "symbol-test.h"

Module-Name ...

Struct test foo;

EXPORT_SYMBOL(foo);

```

static int krt_symbolinit(void)
{
    static void *print_std_type[] = {
        printk(KERN_INFO "i386: Print standard type\n"),
        _func_, foo_a, foo_c,
        return 0;
}

static void __init krt_symbolinit(void)
{
    static void *print_std_type[] = {
        printk(KERN_INFO "i386: Print standard type\n"),
        _func_, foo_a, foo_c,
        return 0;
}

4. Sudo immad .symbol_info
3. f
    41. Version Magic
        It is a magic string present in Linux Kernel and added
        into the .modinfo section of Linux Kernel Modules
        > This is used to verify whether Kernel module was built
        > Compiled for the particular Kernel Version (or) not.
        'VERHATC-STRNG' is generated by Kernel Configuration
        #define VERHATC_STRNG
        "#define VERHATC_STRNG
        "A string" should

```

4 vi Vermagie.

```
#include <linux/vmalloc.h>
```

```
static int printk(KERN_INFO "Hello, world!\n",
```

Principles
return o

R
P
T
S
return 0;
3
Static void.

3
+
Static void ...
three steps ④

Three steps ④
1. Sudo install symbol link

3 f. Verizon Magic

41 Union Magic
It is a magic string present in Linux kernel and added in a module.

into the main section of Linux, I made the following changes:

Compiled for the particular kind of work which has been done.

'VERMATIC-STRINGS' is generated by Kernel Configuration in

11.12.2011 - 11.12.2011

"*U.S.-RELENG*" "4
"A. 4.3-10000P2" *obsolet*

↳ In vmmagic2.c

include <linux/vmmagic.h>

static int (*Save_fn_in_vmmagic_c.c)();

↳ static int (*Save_fn_in_vmmagic_c.c)();

MODULE_INFO(vmmagic, "1234567");

↳ MODULE_INFO(vmmagic,

↳ additionally one statement in vmmagic.c

↳ How we change vmmagic to some 1234567 "value".

↳ \$ sudo insmod ./vmmagic2.ko

↳ \$ rmmod ./vmmagic2.ko

43. What is tainted kernel?
When the kernel is tainted, it means that it is in state that is not supported by kernel community.

→ In addition, some debugging functionality and API calls may be disabled when kernel is tainted.

Kernel may become tainted for any of several reasons, including (but not limited to) the following:

① use of proprietary (or non-English-compatible) kernel module

- this is the most common cause of tainted kernels and usually results from loading Proprietary NVIDIA & AND drivers.

② use of staging drivers, which are part of kernel source code but are not fully tested.

③ use of out-of-tree modules (we are creating) that are not included with Linux Kernel source code.

Contain critical error conditions, such as machine

42. Module license :-
What happens if we don't specify MODULE_LICENSE macro?

→ it will give warning, but it can be loaded successfully.

→ But if we check modinfo for that module, we can't find the license information about it.

↳ \$ modinfo vmmagic
↳ vmmagic: 1234567

Kernel oopen.

```
Export-Symbol Gpl(myAdd);
```

modul. 2.c

45- How to check whether kernel is in tainted state or not

on further ~~modifications~~
print "loading out-of-tree module tainted kernel".
~~taint is tainted~~

→ After that, it will now be
our quarry tainted state by reading

\$ cat /proc/sys/kernel/tainted

that returns 'o', Kond is not tanned.
Indicates the reason why

any other time
watch lecture → for more info

46. what happens when you specify `inner`?

—Almost same as others.

What happens when non GPL kernel module trying to

call an OpenMP Module?

module.c

#include < ->

MODULE-LICENSE(^{mpv})

```
int myadd(int a, int b)
```

— TAKEN IN KICK-STARTED MODE WITH VARIOUS

MODULE-VERSION, macro sets the version of

the module.

- Give the warning before building .ko.
- (i) GPL-incompatible module modules.ko uses GPL-~~only~~
Symbol 'myadd'.
- make V=1 → This gives Verbose Info. (What happens internally
in kernel?)
- Q. How to find out kernel version from .ko?

modifying → gives you magic. [Here we can find].

49. Module Metadatentests ist das Submodul 2 und mit 12 Tests.

MODULE - DESCRIPTION can be short synopsis of what your module is trying to accomplish.

MODULUS AT THE DIRECTION OF THE

We can add all this in our file.

Eg) `modinfo -k Hello.ko`

filename: /home/linux...

version: 1.1.1

license: GPL

author: Linux Trainor

description: Hello world

version: A00f668

depends:

retpoline: *

name: Hello

retpoline: *

name: hello

retpoline: *

name: generic

retpoline: *

name: hello

retpoline: *

name: generic

retpoline: *

name: generic

retpoline: *

name: generic

retpoline: *

name: generic

Retpoline: "Retpoline" was introduced to be a solution to mitigate the risk of Reptile bug. Intel has to be provided patch and we can do it.

Version of generic:

Let's say we provide a ko to client. But client reported that there is a bug. We have released new ko which fixes the bug. But still client explains it doesn't fix. But we are pretty sure that it should fix. So we can check whether client has loaded fix ko(s) or not. We can check "version" of ko file that loaded by using modinfo. Then we can compare with fix file & confirm.

So, MODULE_INFO macro:

cd Linux-5.2.8/

cd include

grep -r 'MODULE_VERSION'.

if # define MODULE_VERSION (_version) MODULE_INFO(version, _version)

grep -r 'MODULE_AUTHOR'.

if # define MODULE_AUTHOR (_version) MODULE_INFO(author, _author)

51. objdump on kernel module

\$ file ./mod_info.ko

objdump

./mod_info.ko: ELF 64-bit LSB relocatable, not stripped.

at build time from modpost script. It's output is in ELF executable Linkable format

It can be changed by 2 options:
General setup → Kernel log buffer

- ① male mandibles → broad jaws spread by default a buffer of max 16-39% bite, little to invoke damage

~~if you use a larger log buffer you have to wait~~

G. # Complex Log-B

Adams - \$ 256000
Int method need to rebuild the entire house
line:

② By using Kernel (

log - bug - len = 44
it is kernel buffer length length to 4 Megabytes

3. DER KERN-INFO-KERN-ALERT

There are eight possible logland strings; defined in the

header & lines / kern - levels. ?

```
#define KERN_SOH "1001 / ASCII string  
#define KERN_SOH-ASCII "1001"
```

81 = 7442 - 318 - 001 - 01100

convole device is not running

55. Default print log level

what happens if we don't specify log level?

5. static modinit(void)

```
printf("-s":Init(n, -fun-),  
return 0;
```

Both doesn't have
log levels.

```
static void exit(void)
{
    printf("...: In exit\n", -fun--);
}
```

When compile, no errors & warnings.

gendo immo./hello.no (05) 8888-
dove -X

old
=> default log level

Kun: Wahr [F 7117-364-] Prod. init: In mit

- Loglevel: level under which messages are output to cat /proc/sys/kernel/printk
- They are associated with following variables:

convole device to hold various liquids and stock was another
example. - Very probably "cannisterns" etc.

default_message_loglevel: priority level that is associated by default with messages for which no priority value is specified. (default with messages from min level to allow a message to be logged on the console device.)

maximum_console_loglevel: maximum level to be logged on the console device.

maximum_loglevel: maximum level

`$ echo 8 > /proc/sys/kernel/printk`

will change console_loglevel

`$ cat /proc/sys/kernel/printk`

(high) 8 (low) 0

Console loglevel
Console kernel
Console default

56. Konsole on Console

What is Konsole?

Console is a device to which we can write text

for Linux, a console is a device to which we can write text

data & read text data.

By default, Console is the (Text mode) screen & keyboard.

One can switch to that console by pressing `Ctrl+Alt+Fn` where n is 1 to 6.

→ When one boots his pc, kernel prints lots of messages, like "initializing this ...", "initializing that ...". These

all get printed via printk that sends message to the console driver. (normally prints prints in browser why?) (cont'd)
so, why don't kernel message appear other in graphic mode?

In Linux, graphics mode is implemented not inside the kernel (& thus it cannot print message in graphics mode), but as a usermode process called X. (or sometimes X server).

→ Every program that wants to 'say', display a window, sends a message via TIPC to X server and says it how it (X-server) should draw the window.

→ Of course, this message passing is implemented in shared library, so from the application writing point of view, it's just a call to function that displays the window.

→ Xterm is one of many graphical application (Konsole and gnome-terminal are other two well known programs that emulate a terminal).

→ PS - \$ | grep 'X' > logs & viewer
gdm. user/birdXwayland:

(3x) X-server → implemented in user space not in Kernel.

↑
very simple understanding

→ X term parses all key strokes to the shell it runs and
→ X term parses all key strokes to the shell it runs and
the generated text is drawn graphically using the

- ① Selected font (by using some functions in shared library that send message to X screen & tell it to draw the characters)

→ kernel messages are not output from shell, so X term does not receive them, and therefore does not print (con... draw) them

To console:

Switch to CTRL+ALT+F1 (or any other showing alt and b)

To return to GUI:

CTRL+ALT+F2

- 5 → print → "welcome" string
ask user → "welcome" string
→ try to load any module
- 6 → check driver
→ check if cat /proc/sys/kernel/printk
of h 1 - 7 → give error about "not known" (or) known wrong
completely
→ if sudo dmesg -n 5 → this also gives known wrong
→ if sudo dmesg -n 100
→ drivers → shows exit message

5.7. More printk macros:

There are other fun that we can use in place of printk() in order to have more efficient message and compact & readable code.

```
#define pr_err(fmt, ...) \
    printk(KERN_EMERG pr_fmt(fmt), #__VA_ARGS__)\n\n#define pr_info(fmt, ...) \
    printk(KERN_INFO pr_fmt(fmt), #__VA_ARGS__)\n\n#define pr_debug(fmt, ...) \
    printk(KERN_DEBUG pr_fmt(fmt), #__VA_ARGS__)\n\n#define pr_notice(fmt, ...) \
    printk(KERN_NOTICE pr_fmt(fmt), #__VA_ARGS__)\n\n#define pr_warning(fmt, ...) \
    printk(KERN_WARNING pr_fmt(fmt), #__VA_ARGS__)\n\n#define pr_err(format, args...) \
    printk(KERN_EMERG format, ##args)\n\n#define pr_info(format, args...) \
    printk(KERN_INFO format, ##args)\n\n#define pr_debug(format, args...) \
    printk(KERN_DEBUG format, ##args)\n\n#define pr_notice(format, args...) \
    printk(KERN_NOTICE format, ##args)\n\n#define pr_warning(format, args...) \
    printk(KERN_WARNING format, ##args)
```

So instead of `printk(KERN_INFO, "....")`, we can use directly `pr_info("....")`. Likewise we can use macro directly.

5.8 Enable pre-debug message :-

Let's say

```
static int hello_init(void)
{
    printk(KERN_DEBUG "1: In init\n", -func-);
}
```

```
pre_debug("1: In init2\n", -func-);
}
}
```

In above only `printk` will print, but `pre_debug` is not printing.

By default, it is not enabled.

To enable:

(1)

In makefile,

obj-m:=hello.o

obj-y:=-DDDEBUG

or flags-y:=-DDDEBUG

(now, make the makefile, build → pre-debug will print)

(59. lkm which prints floating point number)

g.: static int test_hello_init(void)

{ float f = 2.56;

printf("4.s: In init, value of f in %f\n", f); }

return;

}

makefile -

5) Sudo command

3) dmesg → see error

(14) printf will not print floating numbers.

(15) printf will work in printf not in printf.

→ format specifier will work in printf.

60. No use of floating point in kernel.

- when a user-space process uses floating-point instructions, kernel manages the transition from integer to floating point mode.

Why is floating point mode off?

→ Many programs don't use floating point (or) don't use it on any given time slice, and saving FPU registers & other FPU's taken time; therefore an OS kernel may simply turn off FPU off.

process, no state to save and restore, & therefore faster Context-Switching.

If a program attempts an FPU op,

→ the program will trap into kernel.

Kernel will turn FPU on,

restore any saved state that may already exist and then return to re-execute FPU op.

At context switch time, it knows to actually go through the state save logic. (And then it may turn FPU off again.)

The reason that kernel doesn't particularly need FPU ops and also needs to run on architectures without an FPU at all.

Therefore, it simply avoids the complexity & runtime required to manage its own FPU context by not doing ops for which there are always other software solutions.

6. Limiting printk messages - printk_rate_limit!

- o printk is used in kernel printing the kernel prints.
- o useful while debugging the kernel prints.
- o But kernel log being ~~too much~~ big, is generated & unnecessary.
- o Logging at time can lead to looking relevant messages.

7. Then kernel provides a function, printk_ratelimit to restrict the logging using which we can set a limit on the no. of prints that we want our program to do.

The limit on no. of prints is set in file /proc/sys/kernel/printk

- ratelimit_burst

\$ cat /proc/sys/kernel/printk_ratelimit_burst

10

5 The printk_ratelimit function will allow 10 prints before it starts blocking further prints.

The printk_rate_limit for returns 1 as long as no. of prints

do not exceed the limit. Once the limit is reached, it

Start returning 0.

Thus it can be used as a condition for an "if" statement to decide whether to print a message (or) not.

printk will be enabled again after time interval in sec mentioned in file /proc/sys/kernel/printk_ratelimit

```
#include <linux/kernel.h>
#include <linux/module.h>
#include <linux/printk.h>
```

```
#include <linux/delay.h>
```

```
#include <asm/types.h>
```

```
static int hello_init(void)
```

```
{ int i;
```

```
for(i=0; i<200; i++) {
```

```
    if(prin
```

generally it needs sleep to next printk to next printk for now we don't mind anything to do else loop runs for 10 times

return 1 for first 10 times

then return 0.

else {

pr_info("printing %d\n", i+1);

pr_info("Sleeping for 5 seconds\n");

}

static void hello_exit(void)

pr_info("Done\n");

}

0) printing 1
printing 2

printing 10

sleeping for 5 seconds } 10 times

sleeping for 5 seconds

lets try the same with below changes

for 2 ("Sleeping for 5 seconds");

pre_info ("Sleeping for 5 seconds");
nuseep (5000);

printk-once is fairly trivial - no matter how often if we call it, it prints once & never again.

3) #include <

static int hello_init(void) {

int i; for(i=0; i<20; i++)

→ printk-once (KERN_INFO "printing %d\n");

return 0;

3) (1) initialization

if we change

for(i=0; i<20; i++)

→ printk-once (KERN_INFO "printing %d\n"); → only one

3) printk-once (KERN_INFO "printing %d\n"); → therefore
printk-once (KERN_INFO "printing %d\n"); → print
return 0;

)

(iv) This is for line by line, not for the whole module.

④ if we use, printk-once (KERN_WARN, "...") will error. we can't use KERN_WARN with once.

6.3. Avoiding default newline behaviour of print

(
E:
static int hello_init(void)
{
 printk(KERN_INFO "Hello", __func__);
 printk(KERN_INFO "World\n");
 return 0;
}

(3)
off: hello_init: Hello
World

Why "World" doesn't print with Hello?
Kernel didn't add "\n" to Hello?

E:
(ii) Default behaviour of print is adding newlines.

E:
How to print a message in one single line in Linux kernel?

To prevent a new line from being started, we

KERN_CONT → kernel continues

(i.e.) KERN_CONT → kernel continues

Static . → Not Log Level (7) → No. of bytes to dump

printk(KERN_CONT "Hello", __func__);

printk(KERN_CONT "World\n");

printk(KERN_CONT "World\n");

6.4. Printing hex dump - print_hex_dump

Void print_hex_dump (const char * level, const char * prefix-str,
int prefix-type, int rowsz, int groupsize, const void *buf,
size_t len, bool asciin);

Given a buffer print_hex_dump prints a hex+ASCII dump to
the kernel log at specified kernel log level with an optional
leading prefix.

level: kernel log level (e.g. KERN_DEBUG)

prefix-str: string to prefix each line with;

prefix-type: Controls whether prefix of an offset, address or raw
is printed (1: DUMP_PREFIX_OFFSET, \ "DUMP_PREFIX_ADDRESS"

\ "DUMP_PREFIX_NONE")

groupsize: no. of bytes to print per line; must be 16(0x32).

bytes: data blob to dump

E: #include <linux/print.h> → includes all the necessary headers
and defines ASCII after hex of (as far as you can see)

6.5. Hex dump printing function

Static int test_init(void)

{
char buf[10] = "Hello world. Linux is the best os";

char buf[10] = "Hello world. Linux is the best os";

char buf[10] = "Hello world. Linux is the best os";

char buf[10] = "Hello world. Linux is the best os";

char buf[10] = "Hello world. Linux is the best os";

char buf[10] = "Hello world. Linux is the best os";

char buf[10] = "Hello world. Linux is the best os";

char buf[10] = "Hello world. Linux is the best os";

char buf[10] = "Hello world. Linux is the best os";

char buf[10] = "Hello world. Linux is the best os";

char buf[10] = "Hello world. Linux is the best os";

char buf[10] = "Hello world. Linux is the best os";

char buf[10] = "Hello world. Linux is the best os";

char buf[10] = "Hello world. Linux is the best os";

char buf[10] = "Hello world. Linux is the best os";

char buf[10] = "Hello world. Linux is the best os";

char buf[10] = "Hello world. Linux is the best os";

char buf[10] = "Hello world. Linux is the best os";

char buf[10] = "Hello world. Linux is the best os";

char buf[10] = "Hello world. Linux is the best os";

char buf[10] = "Hello world. Linux is the best os";

6. Dynamic Debugging:-

\$ cd /boot

\$ boot & uname -r

5.0.0-23-generic

boot \$ cat config-5.0.0-23-generic | grep CONFIG_DYNAMIC_DEBUG

CONFIG_DYNAMIC_DEBUG=y

→ Using pre-debug() globally will cause a tremendous amount of logging so it is not very practical. To make the debug level more manageable, dynamic debugging was introduced.

→ We can activate it with CONFIG_DYNAMIC_DEBUG

→ DYNAMIC_DEBUG is designed to allow you to dynamically probe/disk kernel code to obtain additional kernel information.

→ Currently, if CONFIG_DYNAMIC_DEBUG is set, then all pre-debug()

/dev-dbg() and print_hex_dump_debug() calls can be dynamically enabled per-call [into module(s) file(s) folder].

Kindly watch lecture for more info.

68. What happens if I try to load non-.ko file with insmod?

touch hello.ko

chmod +x hello.ko

file hello.ko

insmod hello.ko: Empty

rmmod hello.ko → Err! Couldn't get module from 'hello': Invalid argument

\$ sudo insmod ./hello.ko

Output: insmod: ERROR: Could not insert module ./hello.ko

parameters

6. Strace on insmod command:

↳ strace -s

\$ strace -o /tmp/strace.out insmod ./hello.ko

Shows lot -

(last - "finit-module ()")

↳ system calls:

↳ strace -o /tmp/strace.out insmod ./hello.ko

↳ 7. Find out name of module from .ko

↳ \$ vi modname-test.c

#include <linux/module.h>

char *modname = STRINGIFY(KBUILD_MODNAME);

static int init = (void)

printf("Hello %s: In init\n", modname);

printk("Hello %s: Module Name: %s loaded\n", modname);

↳ Output: In init "Hello modname-test" loaded

↳ How it works?

\$ make V=1 KBUILD_MODNAME

We can find it on terminal.

7. How to dump kernel stack?

Calling dump_stack() will cause a stack trace to be printed at that point.

↳ \$ strace ./myinit(void)

↳ pre_info("dump_stack myinit\n");

↳ dump_stack();

↳ pr_info("dump_stack after\n");

↳ return 0;

↳ kernel will do "dump_stack" like: will have some

↳ dump_stack() function which will print stack trace.

8. Kernel panic:-

↳ kernel panic in an error in kernel code.

→ On kernel panic, kernel stop running immediately to avoid the

data loss or other damage.

The reason to stop running is to protect your computer.

for example, if the module which is controlling the fan fails to load, kernel generates a panic & freezes the system.

What are reasons for kernel panic?

→ HW or SW issue (e.g. unable to start init process)

→ Bug in kernel driver

→ defective (or) incompatible RAM

what happens when panic, it calls panic() function which dumps some debug information & depending on configuration which reboot the system.

Configure kernel to reboot on kernel panic

By default, kernel will not reboot on kernel panic. There are two ways, by which you can instruct kernel to reboot.

1. Kernel Command Line: Add "panic=N" to the kernel command line, for the kernel to reboot after N seconds. (Don't explain)
2. proc file system: echo N > /proc/sys/kernel/panic for kernel to reboot after N seconds on reboot. Note this setting is not persistent on reboot. (i.e. kernel restarts after reboot)

2nd method:

\$ cat /proc/sys/kernel/panic

0 or -

\$ echo '5' > /proc/sys/kernel/panic when we do this, it depends

If permission denied, then a message kernel boot failed

\$ sudo -s

echo '5' > /proc/sys/kernel/panic

7. Kernel Panic Example :-

Static int test_panic_init(void)

panic("KERNEL-INFO").S: in init\0, _func_;

④ panic ("Hello Kernel I am causing the panic\0");
return 0;

↑ It got S instead
of panic

Once we insert this module, Linux gets restarted.
After restart, we can check value, it will be 0.
→ Yes, it works ④ with 10 s.

¶ What is oops?

An oops is similar to segfault in user space. Kernel throws oops message when an exception such as accessing invalid memory location happens in the kernel code.

- ① kills the offending process
- ② Prints the information which can help developer to debug
- ③ Continuous execution.

Note: After oops, system cannot be booted further as the some of the locks (or) structures may not be cleaned up.

An oops message contains the following information:

→ processor status

→ Content of CPU registers at time of exception.

→ stack tree

→ Call trace

75. oops example

```
static int --
```

```
+ (int*) (0x12) = 'a';
```

```
}
```

```
if
```

```
else if
```

BUG_ON(condition)

(ex) → if this condition true, then it call if (condition) → BUG() i.e we understand that error is done. Problem happened in some code.

BUG()

→ grep . in drivers folder for BUG_ON (gave 5000 plus and for WARN_ON gave more than 6000 count)

What does BUG() macro do?

→ prints the contents of registers, → prints stack tree

→ current process dies

What does WARN() macro do?

→ print the contents of registers → prints stack trace.

Si

static int ...

{

→ should not be directly called, just to simulate.

BUG();

} → If segmentation fault (Core dumped)

off segmentation fault (Core dumped)

\$ sudo dmesg

\$ rmmod -r module_name

\$ modprobe module_name

75. What is BUG & example?

The most frequently found macros in linux device drivers is BUG-ON/BUG and WARN-ON

Please note that after loading this module using insmod, you cannot unload this. If you try to call rmmod, you will get "Module is un" error.

78. How to define preprocessing symbol in Makefile?

(i) Q Static int test_hello(void)

Static int test_hello(void) we enable & tell to makefile?

? #ifdef DEBUG → #if 0

#ifdef DEBUG → #if 0 we enable & tell to makefile

#ifdef DEBUG → #if 0 we enable & tell to makefile

#endif

return 0;

} static void test_hello_exit(void)

? --

? --

makefile

[obj-m:=symbol.o]

#Kbuild understands a make variable named

CFLAGS_modulename_o to add specific flags when compiling

this unit

{ CFLAGS_symbol_o := -DDEBUG }

#This will be applied to all of the source files compiled for

your module with Makefile

#CFLAGS_y := -DDEBUG

all:

clean: ---

80. How to find out how many CPU's are present from user space & kernel space?

cat /proc/cpuinfo

cat /proc/cpuinfo | grep processor

\$ cat /proc/cpuinfo | grep processor | wc -l

From kernel space:

cat /proc/stat

static int

cpu_info("number of online CPUs is %d\n", num_online_cpus);

return 0;

}

① makefile

② sudo insmod

③ check driver.

81. Process representation in Linux kernel and process state

Linux kernel internally refers processes as tasks.

Kernel stores the list of processes in circular doubly linked

list called the task list.

→ Each task/ process is represented in kernel with struct

task_struct (defined in <linux/sched.h>).

→ Thus data structure (task_struct) is huge (~7 kilobytes)

Containing all the information about a specific process.

Let's write a module/driver which reads the circular linked list & prints the following information for us:

list & print the following information for us:

processname \rightarrow process ID \rightarrow process state.

Before that, we should know what are different states approach.

Can be:

TASK_RUNNING(p): process is either currently running or on a

runqueue waiting to run.

TASK_INTERRUPTIBLE(s): process is sleeping / blocked. Can be triggered by a signal.

TASK_UNINTERRUPTIBLE(d): similar to TASK_INTERRUPTIBLE, but

cannot be triggered by a signal.

does not wakeup on a signal.

TASK_STOPPED(r): process execution has stopped. This

happens when task receives SIGSTOP, SIGSTP, SIGTTIN (or)

SIGTSTP signal if it receives any signal while it is

being debugged.

We can find the status using ps -A command.

This ps -A \Rightarrow taken directly from fileproc system, does not

contain buffer.

8.2. LKM example demonstrating process name, process id and process's:

(i) Now, we don't use proc file system to get. we will access Circular double linked list.

print_task.c

```
#include <linux/module.h>
```

```
#include <linux/sched.h>
```

```
char buffer[256];
```

```
char *get_task_state(long state);
```

```
switch(state){
```

```
case TASK_RUNNING:
```

```
    return "TASK_RUNNING";
```

```
case TASK_INTERRUPTIBLE:
```

```
    return "TASK_INTERRUPTIBLE";
```

```
case TASK_UNINTERRUPTIBLE:
```

```
    return "TASK_UNINTERRUPTIBLE";
```

```
case --TASK_STOPPED:
```

```
    return "--TASK_STOPPED";
```

```
case --TASK_TRACED:
```

```
    return "--TASK_TRACED";
```

default :

8 `sprinf(buffer, "unknown type: %ld\n", state);`

gather buffer;

3

static int test_init(void)

8 struct task_struct *task_list;

unsigned int process_count=0; state code state -> func

8 pr_info("-%s: In init\n", func);

/* define for each proc (p) \-----definition-----*/

8 p=&init_task; (p=next_task(p))!=&init_task;`

4 /*

for-each-process(task_list)

5 /*

pr_info("procnum: %st pid: [%d]\t state: %s\n",

task_list->comm, task_list->pid,

task_list->state);`

get_task_state(task_list->state));`

process_count++;

4 pointer

4 sudo command

4 damage file

8.3 LKM example demonstrating Current Macro

While writing a module if we want to get information about the current process that is running in kernel, we need to read the "task_struct" of the corresponding process.

The kernel provides a easy way to do this by providing a macro by the name "current", which always return a pointer

to the "task_struct" of the current executing process.

This macro has to be implemented by each architecture.

Some architecture stores this in a register, some stores

them in bottom of kernel stack of process.

9 #include <linux/init.h>

#include <linux/module.h>

#include <linux/sched.h>

#include <asm/current.h>

static void current_exit(void)

printk("Current pid: %d, current process: %s,\n", current->pid, current->comm);`

Static int current_init(void)

printk("Current pid : %d, Current procen: %s \n",

current->pid, current->comm);

return 0;

}

Ex. LKM Current macro part - 2: Now only few tasks will be

little modification: printed. (i) from current task
to init_task (ii) again from init_task (iii)

struct ...

for (task= current; task!= &init_task; task= task->parent)

{

pr_info("procen: %s \t pid: [%d] \t state: %s \n",

task->comm, task->pid, get_task_state(task->state));

process_count++;

}

85. LKM which accept pid as argument & print procen & parent

Static unsigned int PID = 1; // default
module_param(PID, uint, 0400);

void print_task(struct task_struct * task)

printk("procen: %s, parent procen is - %s \n", task->comm, task->parent->comm);

}

Static int init_find_task(void)

{ struct task_struct *task = NULL;

for_each_process(task){

if (task->pid == (pid_t) PID){

print_task(task);

}

return 0;

}

\$ sudo insmod ./task.ko

of procen : system, parent procen : swapfile

get, find another pid to check

\$ ps -ef | less (we choose 45)

\$ sudo insmod ./task.ko PID=45

of procen: netns, parent procen: kthread

86. procen memory map:

Struct mm_struct → contains list of procen vmas, page tables, etc...

All info related to procen address space is included in an object called memory descriptor of type mm_struct
available via current->mm

3

What is the use of Kernel Thread?

Kernel Thread helps kernel to perform operations in background.

Struct mm_struct {

* pointer to head of list of memory region objects * /

Struct mm-area_struct *mmap;

Struct mm-area_struct + mmap;

* pointer to root of red-black tree of memory region object +

Struct mm-area_struct + mm-root;

* pointer to last generated memory region object +)

Struct mm-area_struct + mmq - cache;

} ;

Linux implements a memory region by means of an object of

type mm-area-struct .

* Get /proc/(20047)/maps.

Q7. Example for printing predominant memory map.

88. Introduction to Kernel Threads:

↳ Kernel threads are created by kernel itself.

↳ It is not created by running fork() or clone() system calls.

↳ It is implemented in Linux.

↳ It is a kernel thread.

→ Kernel thread is a kernel task running only in kernel mode.

It is not created by running fork() or clone() system calls.

Note:- Kernel - Create only creates the thread but doesn't

run the thread, we need to call wake-up-process() with

Example of Kernel Thread:

1. Ksoftirqd is per CPU kernel runs processing softirqd.

2. Kworker is a kernel thread which processes work queue.

Difference between Kernel Thread & user Thread?

Both Kernel thread & user thread are represented by task-struct.

The main difference is that there is no address space in kernel threads. mm variable of task-struct is set to NULL.

How to Create Kernel Thread?

API for creating kernel thread.

#include <linux/thread.h>

Struct task_struct *kthread_create (int (*thread_fn)(void *data),

void * data, const char name[], ...)

Parameters:

thread_fn → the fn which thread should run.

data → argument for thread fn.

name → printk style format for name of the kernel thread.

Return value → pointer to struct task_struct.

Note:- Kernel - Create only creates the thread but doesn't

run the thread, we need to call wake-up-process() with

return value of kthread_create as an argument to the wake-up-process for thread to run.

Alternative: Linux provides API which creates the kernel thread & calls wake-up-process.

struct task_struct * kthread_run (int (*threadfn)(void *),

void * data, const char name[], ...)

To stop the kernel thread:

int kthread_stop (struct task_struct * k);

Note: If you don't stop the kernel thread in your module_exit

function, you will get oops message.

K-thread_stop is a blocking call; it waits until the fn

executed by thread exits. kthread_stop() flag sets a variable

in task_struct variable which function running in thread

while() should check in each of its loop.

§ int threadfun (void * data) {

§ while (!kthread_should_stop()) {

§ // perform operations here.

§ return 0; }

§ static int my_thread_init (void)

§ my_thread = kthread_create (threadfunction, NULL, "THREAD");

89. Kernel Thread Example - KThread - Create :-

```
#include <linux/init.h>
#include <linux/module.h>
```

```
#include <linux/kdev_t.h>
#include <linux/device.h>
```

```
#include <linux/slab.h> //kmalloc
#include <linux/wait.h> //copy-to/from-user
```

```
#include <linux/thread.h> //task_struct
```

```
#include <linux/sched.h> //include <linux/delay.h>
```

```
static struct task_struct *my_thread;
```

```
int thread_function(void *pv)
```

```
{
```

```
    int i=0;
```

```
    while (!kthread_should_stop ())
```

```
        mSleep(1000);
```

```
    return 0; }
```

```
static int my_thread_init(void)
```

```
    my_thread = kthread_create(threadfunction, NULL, "THREAD");
```


92. Can we have two kernel threads with same name?

yes.

93. What happen if we don't we kthread_stop() in thread function?

kthread_stop(thread_fn); → This is blocking call (i.e.)
it waits until thread_fn is
executed fully. But thread
fn doesn't have kthread_stop(), so it is already
executed before we call / remove module [module_exit → all
kthread_stop].

→ It creates oops in Kernel / warning.

94. What happen if we don't call kthread_stop() in module-with-oops will happen.

95. Print procnum_id in kernel module?

When we have multiple processors present in system, and
want to find out on which processor your driver code is
running, we smp_processor_id().

add printk("smp-processor-id %d\n", smp_processor_id());
in thread fn, module_init, module_exit fn.
we can see, module_init run in one processor, thread_fn
is running & changing in procnum & module_exit in another pr.

96. Linux thread example of race conditions:

Ex: 2 threads,

static int in_thread(void *data)

{ int i=0;

printf("Thread %d started\n"); } program

for(i=0; i<loopcount; i++) {

printf("Thread completed %d\n"); }

counter++; }

return 0;

④ If we pass 1. i.e. 900000

⑤ If we pass 1. i.e. 900000

⑥ Thread started { 1st

Thread completed } 1st

Thread completed } 2nd

Thread completed } 3rd

Thread completed } 4th

Thread completed } 5th

Thread completed } 6th

Thread completed } 7th

Thread completed } 8th

Thread completed } 9th

Thread completed } 10th

Thread completed } 11th

Thread completed } 12th

Thread completed } 13th

Thread completed } 14th

Thread completed } 15th

⑦ Linux version - code review:

we have seen function arguments change b/w kernel
versions for example "alloc-OK".

Before 5.0 Linux version:

int alloc_ok(int type, void *addr, unsigned long size);

After 5.0:

int alloc_ok(void *addr, unsigned long size);

If you want to support multiple kernel versions, you will find yourself having to code conditional compilation directives, the way to do this is to compare macro `Linux-VERSION_CODE` to the macro `KERNEL_VERSION`.

Linux-VERSION_CODE: This macro expands to the binary representation of the kernel version. One byte for each part of the version number. Eg $5.0.0 = 0x500000 = 327680$.

Header file: `<linux/version.h>`
from kernel top level makefile.

`Linux-VERSION_CODE = $(VERSION) + 65536 * $(PATCHLEVEL) + 256`

+ \$(SUBLEVEL)

$$\sum 5.0.0 = 5 * 65536 + 0 * 256 + 0 = 327680$$

If we open "top level makefile" in kernel source code, we can see if we open "top level makefile" in kernel source code, we can see `$(lib/modules)/uname -r` build command it uses `$(VERSION) + 65536 * $(PATCHLEVEL) + 256` for generating uname. It does not generate uname for opening anyone.

```
#include <linux/version.h>
static int test_init(void)
{
    pr_info("Linux Version - Code: %u\n", Linux-VERSION_CODE);
```

Kernel Version Macro: #include <linux/version.h>

2 static int test_init(void)

3 pr_info("Kernel Version for 5.0.0 is - %u\n", KERNEL_VERSION(5,0,0));

This macro is used to build an integer code from individual numbers that build up a version number.

`#define KERNEL_VERSION(a,b,c) (((a) << 16) + ((b) << 8) + (c))`

Header file: `<linux/version.h>`

99. UTS_RELEASE

This macro expands to a string denoting the version of the kernel tree.

Header file: `#include <generated/utsrelease.h>`

If `#include <generated/utsrelease.h>`

`static int test_init(void)`

2 pr_info("UTS Release: %s\n", UTS_RELEASE);

```
#include <linux/version.h>
static int test_init(void)
{
    pr_info("UTS Release: %s\n", UTS_RELEASE);
```

3

100. Lined Kernel Module example supporting multiple versions:

```
static int lastInit(void)  
{
```

162. — This macro informs the compiler to place this function into this special section (.init.text) and this memory is discarded at the function call.

Comment in kernel Log
VERNEL-VERSION (2,6,10)

```
#include <linux/version.h>
printk(KERN_INFO "KERNEL VERSION: Hello old kernel v5.1\n",
      __FILE__);
```

VIS-KELLOGG

#define LINUX_VERSION_CODE 0x01000000

```
printk(KERN_INFO "KERNEL VERSION %s",  
      VTS_RELEASE);
```

PrintK(KERN_INFO) kernel version: NetBSD/nodmide Rev: 1.5

I had the module without --init macro

Only we know
I get / back / ~~the~~ return / your hello →

1811-1812
1812-1813
1813-1814
1814-1815

We can see next-hello-int in many J

Once we load the module, [with `--init` macro]

1-t/mai | Kalkmann | 8sep 2010

Quando la prima volta

We can't see him - hello - isn't in memory, because

fr. & called in completed, Knewl (group the men)

that did not have an impression or always had it

```
make -C /lib/modules/`uname -r` build N=$(pwd) hello.o
```

卷之三

103. Can we use —init macro for built-in module?

- 1) a) Loadable b) Built-in
- Modules can be a) Loadable b) all functions flagged with static for all functions flagged with static kernel looks for all functions flagged with static & arranges them so that kernel build system looks for all pieces of kernel & arrange them so that kernel build system looks for all pieces of memory.

—init, across all pieces of memory, they will all be in same block of memory. Then when kernel boots, it can free that one block of memory.

- When kernel boots & see something like all at once
- When you boot your kernel & see something like 236K freed, this is precisely when you boot your kernel memory.

'Freeing unused' what kernel is doing.

↳ `lspci -v` | grep 'freeing unused'

↳ `lsm -exit macro` → —exit macro defined in `libkmod/init.h`

↳ `lsm —exit` → —exit is a macro defined in `libkmod/init.h`

What is benefit?

When your kernel module is configured as built-in instead of loadable, there is no need of exit function.

↳ Adding —exit macro causes the compiler to discard this function when module is configured as built-in. This saves on code space as the exit functions won't be called.

For loadable modules, this macro doesn't make any effect.

106. —initdata & —exitdata macros

- 1) —initdata & —exitdata in for functions & —initdata / —exitdata

→ —init / —exit in for functions & —initdata / —exitdata in for variables.

↳ static char `buf[10]` —initdata = "LWL";

↳ static int `init_text-hello-init(void)`

2) pre-init ("LSI: In init: %s\n", —func--, buf);

return 0;

→ After inserted module, this will also not in RAM.

↳ How do you list built-in modules?

↳ Kernel modules are pieces of code that can be loaded to kernel upon demand.

↳ It can be configured in 2 ways:

1. Built-in (part of kernel) : Always present

2. Loadable : can be loaded/unloaded as per the need.

↳ lsmod command lists all modules that are built into kernel.

↳ cd /lib/modules/`uname -r`/modules/built-in

↳ vi modules.builtin

↳ Shows a blank file with content `built-in`.

108. How to load modules automatically?

① modprobe -r hello

② modprobe hello

③ sudo insmod hello

Creating symbol

④ ls -l

⑤ Sudo depmod -a

⑥ Sudo modprobe hello

⑦ Sudo modprobe -r hello

⑧ Sudo vi /etc/modules

⑨ add ~~hello~~ hello1

⑩ reboot

⑪ lsmod | grep hello

⑫ lsmod | grep el000

⑬ Why do we need to blacklist kernel module?

Blacklisting is a mechanism to prevent KM from loading.

Why do we need to blacklist kernel module?

→ Associated hardware is not needed.

→ If loading that module could cause problem.

Ex. There may be two KMs to control same piece of HW & loading them together would result in conflict.

How to blacklist?

Create a .conf file inside /etc/modprobe.d/ and append a line for each module you want to blacklist using blacklist keyword.

blacklist <module name>

Finally run, the following commands to rebuild initramfs & sudo update-initramfs -u.

\$ cd /etc/modprobe.d/
\$ lsmod | grep el000 > el000.conf [we are going to blacklist]

\$ rmmod el000

\$ sudo -s
\$ lsmod | grep el000

\$ touch /etc/initramfs-tools/modules

\$ vi /etc/initramfs-tools/modules

In this file,

blacklist el000

\$ lsmod | grep el000

\$ modinfo el000 | grep el000

\$ sudo update-initramfs -u
\$ reboot

No space

Kernel Command Line:

10. Blacklisting using `kernel` modules from the boot loader.
We can also blacklist modules, module1, module2, ...

simply add module-blame
to your bootloader's kernel line
to your broken module make it
Note: this can be very useful if a broken system.

impossible to book

g. 15mm | gryp 600
off 1000
g end vi / et / defa

inside that

GRUB-CMPLXUE LINE = finaplix module-blcklist = e1000
priority = Critical. isolaten = en-us 1)

part added

112. 2019

A group - or 'options'

ofp: - - : options b687x index=2

`sysctl` can be used to print values of files present in sysfs module's sysfs directory.

III. Tuning parameters to 4km loaded using malprobe

Setting module options:

To pass op ions to modules which are automatically loaded using read probe.

↳ Create a file in /etc/modprobe.d/ with .conf extension
syntax: options module-name parameter-name = parameter-value

Note: Configuration files in the directory can have any name, given that they end with the .conf extension.

Linux booting

Bios → MBR → GRUB → Kernel → Init → Pseudo

1. Bios :-
i. stands for Basic Input / output system.
ii. performs some system integrity checks.
 2. performs some system integrity checks.
 3. Searches, loads & executes the boot loader program.
 4. It looks for bootloader in floppy, cd-rom (or) hard drive.
we can press a key (F12 or F2) but it depends on system
during Bios startup to change the boot sequence.
 5. Once the bootloader is detected & loaded into the memory,
 6. Bios gives the control to it.
- booting starts with :-
1. MBR :-
i. MBR stands for Master Boot Record.
 2. It is located in 1st sector of the bootable disk. Typically /dev/hda (or) /dev/sda.
- it contains the below things which have size 512 bytes in size. This has 3 components
1. primary boot loader info in 1st 446 bytes
 2. partition table info in next 64 bytes
 - 3.) mbr validation check in last 2 bytes.
4. It contains information about GRUB
5. So in simple terms MBR loads & executes GRUB boot loader.

In my PC:

cd /dev/

ls -l /dev/sda

off -k sda

-k sda1

ls -l /sda1

3. GRUB:-

1. It stands for Grand Unified Bootloader

2. If you have multiple kernel images installed on your system,

you can choose which one to be executed.

3. GRUB displays a splash screen wait for few seconds, if you

don't enter anything it loads the default kernel image

as specified in the grub configuration file.

4. GRUB has the knowledge of the file system & of devices

5. It contains kernel & initrd images.

cd /boot

ls -l /boot/vmlinuz-2.6.18-163.el5

ls -l /boot/initrd-2.6.18-163.el5

6. Go in simple terms GRUB just loads & executes kernel & initrd

inited image.

4 Kernel:-

→ Mount the root file system as specified in /etc/fstab

"root=

→ Kernel executes /sbin/init program.

→ Since init was the 1st program to be executed by Linux kernel, it has process id (PID) of 1. To see "ps -ef | grep init" & check the PID.

→ initrd stands for initial Ramdisk.

→ Initrd is used by kernel as temporary root file system while kernel is booted and real root file system is mounted. It also contains necessary drivers compiled inside which helps it to access the hard drive partitions & other hardware.

5. init:-

→ cd /etc

init file was there.

→ Init identifies the default initlevel & used that to load all appropriate programs.

→ Typically we would set default run level to either 3 (or) 5-6 where we moved away from sysV init to upstart instead.

6. Runlevel program:-

→ When Linux system is booting up, you might say "starting sendmail...ok". These are runlevel programs executed from the run level directory as defined by your run level.

netstat → status of network

↳ netstat -nr → shows kernel IP routing table

↳ netstat -i → shows kernel interface table.

↳ netstat -a → shows active internet connection (servers & established)

↳ shows with domain name.

3.

↳ netstat -tan → same as above command. one difference

is that shows with IP address

host:

When we search for www.google.com → first IP address is given to computer, if it is found. It will run by local net in internet. Then it checks DNS server in internet. That's it.

Let's try

go to any webpage like Indianbank & get the IP from somewhere

↓
\$ sudo su

↳ netstat -at shows download progress

edit here,

61.246.225.227 ip Indianbank

↳ This should be edited before the below statement.

"# The following lines are desirable for IPv4 (spelled host)

then open browser.

Type "telnet" & press enter.

It will show first time error like no such host & go to advent.

↳ now Indianbank page open.

Network tracing:

\$ sudo apt install traceroute

\$ traceroute www.google.com

it shows each & every step to connect google.com

Network map for tracking activity on network:

\$ sudo install nmap

\$ nmap -v "ip-owner"

↳ \$ nmap -v 192.168.43.78

To find range of IP, Connected to device

\$ nmap 192.168.43.78,1-100

otherwise

\$ nmap 192.168.43.78

op1 =

~~Neas~~ for
lot of info.

1 nmap -A 192.0.168.43 -T8

卷之三

with os also

112

Auditorium

- ① Go to git hub, create one repository, take that link
 - ② Go to Linux terminal, create one directory.
 - ③ \$ cd git-hub
 - ④ \$ cd git-hub
 - ⑤ \$ git init → Contains empty .git folder
 - ⑥ \$ vi list.txt
 - ⑦ \$ git add list.txt
 - ⑧ \$ git commit -m "Not added project started"
 - ⑨ \$ git branch -M main
 - ⑩ \$ git remote add origin https://github.com/lviv
 - ⑪ \$ git push origin main → push branch to repo.
 - ⑫ It will ask for username & password.

username → GitHub username

password → not git hub ~~pass~~ password, token has to be generated → click git hub page.

- ⑪ git branch release → creating new branch in local m/c
 - ⑫ git push origin release → push to remote m/c
→ remote name
 - ⑬ you can see changes in git hub.
 - ⑭ ~~git checkout release~~ → do modify
 - ⑮ vi list.txt → do modify
 - ⑯ git add list.txt
 - ⑰ git commit -m "New list" → Commit message
 - ⑱ git push origin release → pushing update release in remote
 - ⑲ U can see "release had recent pushes 11 minutes ago" in git hub → Compare & pull request
 - ⑳ click compare & pull request
 - ㉑ there they ask to merge [base: main] ← [Compare: release]
 - ㉒ click merge. no new change in release is merged to main branch. all in remote.
 - note: in local m/c, release has only new change, main is old list. [Based on our notes]
 - ㉓ In remote, after merge, it will ask "say to delete release"

in local m/c .
it will be available in delete release in local m/c.

it will be available to delete release in local m/c.

it will be available to delete release in local m/c.

it will be available to delete release in local m/c.

it will be available to delete release in local m/c.

it will be available to delete release in local m/c.

it will be available to delete release in local m/c.

it will be available to delete release in local m/c.

it will be available to delete release in local m/c.

it will be available to delete release in local m/c.

it will be available to delete release in local m/c.

it will be available to delete release in local m/c.

it will be available to delete release in local m/c.

it will be available to delete release in local m/c.

it will be available to delete release in local m/c.

it will be available to delete release in local m/c.

it will be available to delete release in local m/c.

it will be available to delete release in local m/c.

it will be available to delete release in local m/c.

③ \$ git push origin main -> try to update nothing
on remote m/c.

④ error: failed to push something.

⑤ release: Because main local m/c doesn't have "sankar" name "update". But Remote main has the update.

⑥ "sankar" name "update" in local m/c has to be up-to-date with remote before committing new change in local branch to remote branch.

⑦ First committing new change in local branch to remote branch.

⑧ \$ git pull origin main

⑨ \$ cat list.txt -> These Sankar name -> updated.

⑩ \$ now push origin.main to remote .

⑪ \$ git push origin main

⑫ \$ can see nothing.txt in remote .

All work done .

⑬ Delete repository in remote:

⑭ click repo name [Audiosteam] .

⑮ \$ go to settings [code issues ... settings]

⑯ \$ go to Danger zone [Delete this repository]

⑰ \$ git add signing.txt

⑱ \$ git commit -m " signing " -> In local main m/c .

note:- remote main is not updated .